

Homework Notes: I did not work with anyone else on this homework or refer to resources other than the course notes, textbook, and course Piazza page.

Problem 1

A To update the sample of μ_k at iteration $m + 1$, we can sample $\mu_{k,m+1} \sim \text{Unif}(l, u)$ where $l = \min(0, \mu_k - \epsilon)$ and $u = \max(5, \mu_k + \epsilon)$. This prevents us from sampling outside the range defined by the prior. In my samples, I set the ‘step’ $\epsilon = 0.1$.

The MH acceptance probability is the ratio of the likelihood of the new and old samples:

$$p(\text{keep}) = \frac{\mathcal{L}(X|\mu_{m+1})}{\mathcal{L}(X|\mu_m)}.$$

B We chose to perform Metropolis-Hastings rather than a Gibbs sample step because we need to sample μ_k for all k simultaneously. Moreover, we cannot sample from the conditional probability $p(\mu_k|x, z, \Sigma)$, which would be needed for the Gibbs step.

C I ran 500 rounds of burn-in and 3,000 iterations of sampling. Parameters were initialized as:

Listing 1: R Code for Initial Parameters

```
1 # settings
2 BURN = 500
3 M = 3000+BURN
4 STEP = 0.1
5 MIN = 0
6 MAX = 5
7 K = 2
8
9 # initialize
10 pis = mus = matrix(NA, nrow=M, ncol=K)
11 mus[1, ] = sample(X, 2)
12 mus[1,]
13 Z = matrix(NA, nrow=nrow(X), ncol=K)
14 Z[1,] = pis[1, ] = rep(1/K, K)
15 N = x.bar = rep(0, K)
16 z = rep(0, nrow(X))
17 alpha = rep(1, K)
18 S_0 = rep(3, K)
19 v_0 = rep(5, K)
20 v = rep(0, K)
21 S = rep(0, K)
22 Sigma.inv = rgamma(K, shape=v_0, rate=S_0)
23 new.mu = rep(NA, K)
```

The following code runs the burn-in and sampling iterations:

Listing 2: R Code for MH

```

25 for (m in 2:M){
26   # step 1 - simulate proportion vector from Dirichlet
27   pisamp = rdirichlet(1, c(alpha[1]+N[1], alpha[2]+N[2]))
28   pis[m, ] = pisamp[1, ]
29
30   # step 2 - sample latent indicators
31   for (i in 1:nrow(Z)){
32     for (j in 1:ncol(Z)){
33       Z[i, j] = pis[m, j] * dnorm(X[i], mus[m-1, j], 1/Sigma.inv[j])
34     }
35     Z[i, ] = Z[i, ]/sum(Z[i, ])
36
37     # step 3 - calculate new z's based on Z draws
38     z[i] = sample(seq(1:K), size=1, prob=Z[i, ])
39   }
40
41   # step 4 - update mean and variance
42   for (k in 1:K){
43     want = which(z==k)
44     N[k] = length(want)
45     subset = X[want]
46     x.bar[k] = mean(subset)
47     v[k] = v_0[k] + N[k]
48     S[k] = S_0[k] + sum.sq.err(subset, mus[m-1, k]) # test this
49     Sigma.inv[k] = rgamma(1, shape=v[k], rate=S[k])
50
51     lower = max(c(MIN, mus[m-1, k]-STEP))
52     upper = min(c(MAX, mus[m-1, k]+STEP))
53     new.mu[k] = runif(1, lower, upper)
54
55   }
56
57   ll.new.mu = log.lik(X, z, new.mu)
58   ll.old.mu = log.lik(X, z, mus[m-1, ])
59
60   acceptance.prob = ll.new.mu / ll.old.mu
61   acceptance.prob = min(1, acceptance.prob, na.rm=TRUE)
62   if (runif(1,0,1) < acceptance.prob){
63     mus[m, ] = new.mu
64     likelihood.trace[m, 1] = ll.new.mu

```

```

65 } else {
66     mus[m, ] = mus[m-1, ]
67     likelihood.trace[m, 1] = ll.old.mu
68 }
69
70 if(n%%10==0){ print(m) }
71 }

```

D Figure 1 shows the log likelihood trace for three different runs of the sampler using three different (randomly chosen) starting points.




Figure 1: Traceplot of Likelihood for Three Starting Points

E Figure 2 presents histograms of the posterior samples for each mean parameter from a single run (after burn-in). It appears that $K = 2$ is a reasonable fit for the data, with $\mu \approx (1.5, 3)$. Label switching did not occur in these samples.




Figure 2: Histogram of Posterior Samples

F I would use a composite from all three samples, weighted by the likelihood of each.